

```
In [1]: import numpy as np
import pandas as pd
import matplotlib.pyplot as plt

In [4]: df=pd.read_csv('Salary_Data.csv')

In [5]: df.head()

Out[5]:
  YearsExperience  Salary
0              1.1  39343.0
1              1.3  46205.0
2              1.5  37731.0
3              2.0  43525.0
4              2.2  39891.0

In [6]: df.tail()

Out[6]:
  YearsExperience  Salary
25              9.0 105582.0
26              9.5 116969.0
27              9.6 112635.0
28             10.3 122391.0
29             10.5 121872.0

In [7]: df.info()

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 30 entries, 0 to 29
Data columns (total 2 columns):
 #   Column        Non-Null Count  Dtype  
---  --
 0   YearsExperience  30 non-null     float64 
 1   Salary          30 non-null     float64 
dtypes: float64(2)
memory usage: 608.0 bytes

In [9]: df.isnull().sum()

Out[9]:
YearsExperience    0
Salary             0
dtype: int64

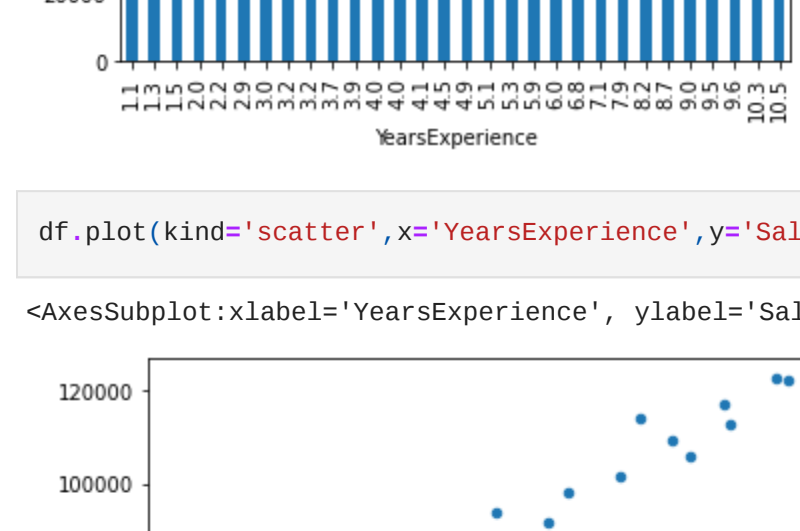
In [13]: df.plot(x='YearsExperience',y='Salary')

Out[13]:
<AxesSubplot: xlabel='YearsExperience'>
```



```
In [14]: df.plot(kind='bar',x='YearsExperience',y='Salary')

Out[14]:
<AxesSubplot: xlabel='YearsExperience'>
```

A bar plot showing the relationship between YearsExperience (x-axis, ranging from 0 to 10) and Salary (y-axis, ranging from 0 to 120,000). The plot shows a clear upward trend, with blue bars representing the 'Salary' data. The bars start at approximately 39,000 for 1.1 years of experience and rise to about 122,000 for 10.3 years of experience.


```
In [15]: df.plot(kind='scatter',x='YearsExperience',y='Salary')

Out[15]:
<AxesSubplot: xlabel='YearsExperience', ylabel='Salary'>
```




```
In [16]: plt.figure(figsize=(10,5))
plt.plot(df.YearsExperience,df.Salary)

Out[16]:
<matplotlib.lines.Line2D at 0x28e2fb35108>
```




```
In [17]: plt.figure(figsize=(10,5))
plt.scatter(df.YearsExperience,df.Salary)

Out[17]:
<matplotlib.collections.PathCollection at 0x28e2fb08c70>
```




```
In [18]: plt.figure(figsize=(10,5))
plt.scatter(df.YearsExperience,df.Salary,color='red')

Out[18]:
<matplotlib.collections.PathCollection at 0x28e2fb28e0>
```



```
In [19]: plt.figure(figsize=(10,5))
plt.scatter(df.YearsExperience,df.Salary,color='green',marker='*')

Out[19]:
<matplotlib.collections.PathCollection at 0x28e30bf640>
```

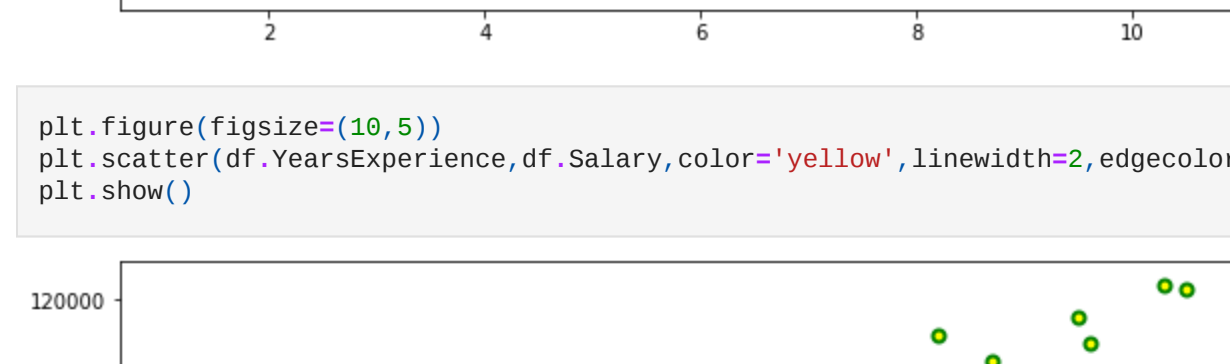


```
In [20]: plt.figure(figsize=(10,5))
plt.scatter(df.YearsExperience,df.Salary,color='red',linewidth=5)

Out[20]:
<matplotlib.collections.PathCollection at 0x28e30c521c0>
```



```
In [21]: plt.figure(figsize=(10,5))
plt.scatter(df.YearsExperience,df.Salary,color='yellow',linewidth=2,edgecolor='green')
plt.show()
```



FROM SCRATCH

```
In [24]: x=df.iloc[:,1:1].values
y=df.iloc[:,1].values

In [25]: x

Out[25]: array([[ 1.1],
 [ 1.3],
 [ 1.5],
 [ 2. ],
 [ 2.2],
 [ 2.9],
 [ 3. ],
 [ 3.2],
 [ 3.2],
 [ 3.7],
 [ 3.9],
 [ 4. ],
 [ 4. ],
 [ 4.1],
 [ 4.5],
 [ 4.8],
 [ 5.1],
 [ 5.1],
 [ 5.9],
 [ 6. ],
 [ 6.5],
 [ 7.1],
 [ 7.1],
 [ 7.9],
 [ 8.2],
 [ 8.7],
 [ 9. ],
 [ 9.5],
 [ 9.6],
 [10.3],
 [10.5]])

In [26]: y

Out[26]: array([[ 39343.,  46205.,  37731.,  43525.,  39891.,  56642.,  60150.,
  64445.,  64445.,  57189.,  63218.,  55794.,  56957.,  57081.,
  61111.,  67938.,  66829.,  63088.,  81383.,  93940.,  91738.,
  98273., 101302., 113812., 109431., 105582., 116969., 112635.,
 122391., 121872.]])

w0=y_mean-w1*x_mean

w1= sum(x-x_mean)*(y-y_mean)/sum(x-x_mean)^2
```

```
In [27]: # Mean X and Mean_Y
mean_x=np.mean(x)
mean_y=np.mean(y)

# Total length of the Dataset
n=len(x)

# Using The formula calculate the w0 and w1
numer=0
denom=0

for i in range(n):
    numer+=x[i]*mean_x*(y[i]-mean_y)
    denom+=x[i]*mean_x**2


w1=numer/denom
w0=mean_y-(w1*mean_x)
print('The Coefficients are',w0,w1)

The Coefficients are [25792.20919867] [9449.96232146]
```

```
In [28]: plt.scatter(x,y,color='m')

# calculate the y_pred
y_pred=w0+w1*x

#plot the regressor line
plt.plot(x,y_pred,color='g')
```



```
In [29]: y

Out[29]: array([[ 39343.,  46205.,  37731.,  43525.,  39891.,  56642.,  60150.,
  64445.,  64445.,  57189.,  63218.,  55794.,  56957.,  57081.,
  61111.,  67938.,  66829.,  63088.,  81383.,  93940.,  91738.,
  98273., 101302., 113812., 109431., 105582., 116969., 112635.,
 122391., 121872.]])

In [30]: y_pred

Out[30]: array([[ 36187.15875227,  38077.15121656,  39967.14368085,  44692.12484158,
  45582.11730987,  53197.09093089,  54142.08710303,  56032.07962732,
  56032.07962732,  60757.06078095,  62647.05325234,  63950.04548449,
  63952.04948449,  64537.04571663,  68317.03064522,  72097.0155738 ,
  73987.00803089,  75877.00605239,  81346.07789525,  82491.9741274 ,
  90051.94398456,  92886.932681 , 100446.90253816, 103281.8912346 ,
 108006.07239533, 110841.86189176, 115566.84225249, 116511.83848464,
 123126.81218966, 125016.80457395]])

In [31]: from sklearn.metrics import r2_score
r2_score(y,y_pred)*100

Out[31]: 95.69566641435085
```

From Using Libraries

```
In [32]: from sklearn.linear_model import LinearRegression

In [33]: model=LinearRegression()

In [34]: model.fit(x,y)

Out[34]: LinearRegression()
```

```
In [35]: y_pred_lib=model.predict(x)

In [36]: y

Out[36]: array([[ 39343.,  46205.,  37731.,  43525.,  39891.,  56642.,  60150.,
  64445.,  64445.,  57189.,  63218.,  55794.,  56957.,  57081.,
  61111.,  67938.,  66829.,  63088.,  81383.,  93940.,  91738.,
  98273., 101302., 113812., 109431., 105582., 116969., 112635.,
 122391., 121872.]])

In [37]: y_pred_lib

Out[37]: array([[ 36187.15875227,  38077.15121656,  39967.14368085,  44692.12484158,
  45582.11730987,  53197.09093089,  54142.08710303,  56032.07962732,
  56032.07962732,  60757.06078095,  62647.05325234,  63950.04548449,
  63952.04948449,  64537.04571663,  68317.03064522,  72097.0155738 ,
  73987.00803089,  75877.00605239,  81346.07789525,  82491.9741274 ,
  90051.94398456,  92886.932681 , 100446.90253816, 103281.8912346 ,
 108006.07239533, 110841.86189176, 115566.84225249, 116511.83848464,
 123126.81218966, 125016.80457395]])

In [38]: plt.scatter(x,y,color='b')

#plot the regressor line
plt.plot(x,y_pred_lib,color='r')
```



```
In [39]: from sklearn.metrics import r2_score
r2_score(y,y_pred_lib)*100

Out[39]: 95.69566641435085

In [40]: x=[[12]]
model.predict(x)

Out[40]: array([139191.74805613])

In [41]: w0+w1*x
25792.20919867* (9449.96232146*12)

Out[41]: 139191.74805619
```

End